

Zora Che

shinechezx@gmail.com • zora-che.com • [Google Scholar](#)

EDUCATION

University of Maryland, College Park Computer Science PhD Student. Advisor: Furong Huang. Vitalik Buterin & Cooperative AI Fellowships. GPA: 4.0/4.0	2023–Present
Boston University — BA Mathematics, Computer Science (Honors), <i>summa cum laude</i> . GPA: 3.96/4.0	2022

SELECTED PUBLICATIONS (FULL LIST ON GOOGLE SCHOLAR)

- **Z. Che**, J. Poveda, A. B. Rodriguez, Y. Y. He, C. B. C. Zhang, Z. Wang, U. M. Sehwag. “Script Kiddie Uplift: Measuring Procedural Misuse Amplification in Large Language Models.” *Spotlight at AI in the Wild Workshop at ICLR*, 2026.
- B. Wei*, **Z. Che***, N. Li, U. M. Sehwag, J. Götting, S. Nedungadi, J. Michael, S. Yue, D. Hendrycks, P. Henderson, Z. Wang, S. Donoughe, M. Mazeika. “Best Practices for Biorisk Evaluations on Open-Weight Bio-Foundation Models.” *NeurIPS Workshop on Biosecurity Safeguards for Generative AI*, 2025.
- **Z. Che***, S. Casper*, A. Satheesh, R. Gandikota, D. Rosati, S. Slocum, L. McKinney, Z. Wu, Z. Cai, B. Chughtai, F. Huang, D. Hadfield-Menell. “Model Tampering Attacks Enable More Rigorous Evaluations of LLM Unlearning.” *TMLR*, 2025. Previously in *Safe GenAI Workshop at NeurIPS*, 2024.
- M. Panaitescu-Liess, **Z. Che**, B. An, Y. Xu, P. Pathmanathan, S. Chakraborty, S. Zhu, T. Goldstein, F. Huang. “Can Watermarking Large Language Models Prevent Copyrighted Text Generation and Hide Training Data?” *AAAI*, 2025. *Best Paper at AdvML-Frontiers at NeurIPS*, 2024.
- B. An, **Z. Che**, M. Ding, F. Huang. “Transferring Fairness under Distribution Shifts via Fair Consistency Regularization.” *NeurIPS*, 2022. Previously in *SRML at ICLR*, 2022.

PRESENTATIONS

“AutoDQM and RPC Monitoring” — CMS Week at CERN	2020
---	------

RESEARCH EXPERIENCE

Fellow, Astra Program (supervised by David Lindner, Google DeepMind) • Research on chain-of-thought monitorability and safety optimization pressure	Jan 2026–Present
PhD Research Intern, Scale AI (supervised by Chen Xing, Udari Sehwag) • Biosecurity and cybersecurity uplift research, hosted by post-training & safety teams	Jun–Sep 2025
Research Scholar, ML Alignment & Theory Scholars (MATS) (supervised by Stephen Casper) • Robust unlearning evaluation and interpretability for LLMs	Jun 2024–Jan 2025
Research Assistant, University of Maryland (supervised by Furong Huang) • LLM safety, online alignment, watermarking, fairness	Sep 2023–Present
NSF REU Scholar, University of Maryland (supervised by Furong Huang) • Fair domain adaptation and adversarial training	Jun 2021–Jun 2022
Institute for Research and Innovation in Software for High Energy Physics Fellow (supervised by Oksana Shadura) • Data infrastructure for high energy physics analysis facility	Jan–Jun 2021
Research Assistant, Boston University HEP Group (supervised by Indara Suarez) • Anomaly detection for CMS collider data quality at CERN	May 2020–Jan 2021

OTHER EXPERIENCE

Data Consultant, Monitoring and Evaluations team, charity: water	Jun–Dec 2019
Editorial Intern, The New York Times	Jul–Aug 2018

HONORS & GRANTS

- Vitalik Buterin PhD Fellowship in AI Existential Safety (2025)
- Cooperative AI PhD Fellowship — 3 years of PhD tuition and stipend (2025)
- Barry M. Goldwater Scholar — awarded to 408 U.S. college students for excellence in research (2021)
- New York Times College Scholar — awarded to 10 graduating seniors in NYC (2018)

SERVICE

- Reviewer: ICLR 2024, NeurIPS SafeGenAI Workshop 2024, NeurIPS 2025 | Volunteer: FAccT 2022, NeurIPS 2022, EAAMO 2023
- AI4ALL Portfolio Project Mentor (2024). Technica Mentor Panelist (2023). Impact Labs Team Lead (2019–2023).

TECHNICAL SKILLS

- Python, Java, C, R; PyTorch, TensorFlow. Methods: RLHF, DPO, finetuning, red-teaming, evaluation design, human studies

ADDITIONAL

- Multimedia artist with exhibitions at Gray Area, MoMA PS1, Chashama Open Studios; writing in Kernel Magazine
- Authorized to work in the U.S. without sponsorship